

# ME 5824: Homework-4

Shaunak Mehta, Ananth Jonavittula

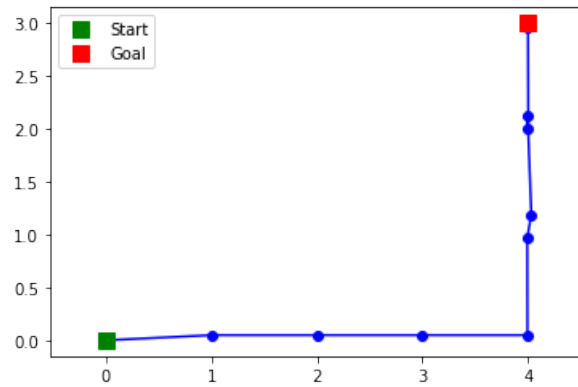
## Q 1

The cost function used for the task is:

$$cost = \sum_s [||g - s|| + s_y - (s_y | s_y < 0.05)]$$

Here  $\sum_s$  sums over the entire trajectory,  $g$  is the goal,  $s$  is the current state in the trajectory,  $s_y$  is the height of the quadcopter, and  $(s_y | s_y < 0.05)$  penalizes the quadcopter if it is too close to the ground.

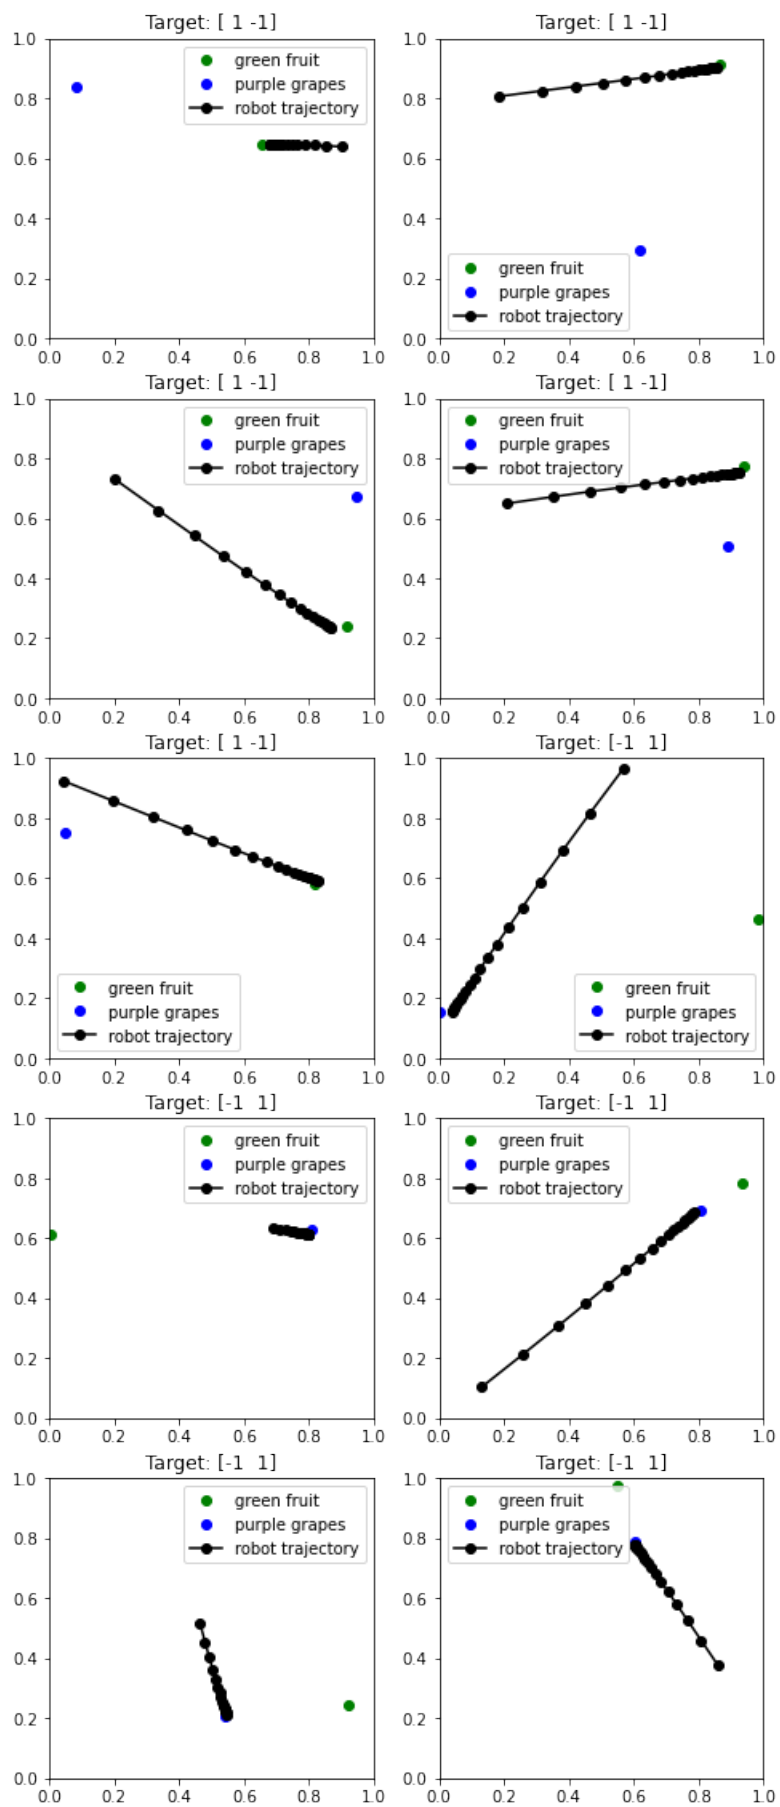
The optimized trajectory obtained at the end of simulation is shown below:



The python script for the same is attached at the end of the file.

## Q 2

Plots can be found on the next page.



The python script for the same is attached at the end of the file.

### Q 3

$$\begin{aligned}
Regret_i &= R(\zeta^*, \theta^*) - R(\zeta^i, \theta^*) \\
Regret_{i+1} &= R(\zeta^*, \theta^*) - R(\zeta^{i+1}, \theta^*) \\
Regret_{i+1} - Regret_i &= R(\zeta^i, \theta^*) - R(\zeta^{i+1}, \theta^*) \\
&= \theta^* F(\zeta^i) - \theta^* F(\zeta^{i+1}) \\
&= \theta^* [F(\zeta^i) - F(\zeta^{i+1})]
\end{aligned}$$

From the hint (assuming that  $F(\zeta)$  and  $\theta$  are unit vectors), we know that  $F(\zeta^i) = \theta^i$

$$\begin{aligned}
Regret_{i+1} - Regret_i &= \theta^* [\theta^i - \theta^{i+1}] \\
&= \theta^* [\theta^i - (\theta^i + \alpha(F(\zeta^*) - F(\zeta^i)))] \\
&= \theta^* [\alpha(F(\zeta^i) - F(\zeta^*))]
\end{aligned}$$

We know that  $\alpha, \theta^* \in [0, \infty)$  and that  $R(\zeta^*, \theta^*) > R(\zeta, \theta^*)$

Therefore,  $F(\zeta^*) > F(\zeta^i)$

Therefore,  $F(\zeta^i) - F(\zeta^*) < 0$

Therefore,  $\theta^* [F(\zeta^i) - F(\zeta^*) < 0]$

Therefore,

$$\begin{aligned}
Regret_{i+1} - Regret_i &= \theta^* [\alpha(F(\zeta^i) - F(\zeta^*))] \\
Regret_{i+1} - Regret_i &\leq 0 \\
Regret_{i+1} &\leq Regret_i
\end{aligned}$$

```
### Link to Github Repo with the same code: ###  
### https://github.com/dylan-losey/me5824.git ###
```

```
import matplotlib.pyplot as plt  
import numpy as np  
from scipy.optimize import minimize, LinearConstraint, NonlinearConstraint
```

```
### Create a class to perform the trajectory optimization ###
```

```
class TrajOpt(object):
```

```
    def __init__(self):
```

```
        # initialize trajectory
```

```
        self.n_waypoints = 10
```

```
        self.n_dof = 2
```

```
        self.home = np.array([0., 0.])
```

```
        self.xi0 = np.zeros((self.n_waypoints, self.n_dof))
```

```
        self.xi0 = self.xi0.reshape(-1)
```

```
        # create start constraint and action constraint
```

```
        self.B = np.zeros((self.n_dof, self.n_dof * self.n_waypoints))
```

```
        for idx in range(self.n_dof):
```

```
            self.B[idx,idx] = 1
```

```
        self.lincon = LinearConstraint(self.B, self.home, self.home)
```

```
        self.nonlincon = NonlinearConstraint(self.nl_function, -1.0, 1.0)
```

```
        # each action cannot move more than 1 unit
```

```
    def nl_function(self, xi):
```

```
        xi = xi.reshape(self.n_waypoints, self.n_dof)
```

```
        actions = xi[1:, :] - xi[:-1, :]
```

```
        return np.linalg.norm(actions, axis=1)
```

```
        # trajectory cost function
```

```
    def trajcost(self, xi):
```

```
        xi = xi.reshape(self.n_waypoints, self.n_dof)
```

```
        cost = 0
```

```
        ### define your cost function here ###
```

```
        ### here is an example encouraging the robot to reach [5, 2] ###
```

```
        for idx in range(self.n_waypoints):
```

```
            cost += np.linalg.norm(np.array([4., 3.]) - xi[idx, :])
```

```
            cost += 1*abs(xi[idx, 1])
```

```
            if xi[idx, 1] < 0.05:
```

```
                cost -= 1*xi[idx, 1]
```

```
        return cost
```

```
        # run the optimizer
```

```
    def optimize(self):
```

```
        res = minimize(self.trajcost, self.xi0, method='SLSQP', constraints={self.lincon, self
```


```
        xi = res.x.reshape(self.n_waypoints, self.n_dof)
```

```
        return xi, res
```

```
### Run the trajectory optimizer ###
```

```
trajopt = TrajOpt()
xi, res = trajopt.optimize()
print(xi)
plt.plot(xi[:,0], xi[:,1], 'bo-')
plt.plot(0, 0, 'gs', markersize=10, label='Start')
plt.plot(4, 3, 'rs', markersize=10, label='Goal')
plt.legend()
plt.axis("equal")
plt.show()
```

```
[[0.          0.          ]
 [0.99879898  0.0489958  ]
 [1.99879891  0.04901323]
 [2.99879884  0.04899491]
 [3.99387065  0.04898458]
 [3.99188275  0.97197579]
 [4.03034504  1.18648881]
 [3.99522503  2.13092249]
 [4.00148332  1.99750678]
 [3.99855789  2.96269541]]
```



# behavior\_cloning

February 26, 2022

```
[ ]: ### Link to Github Repo with the same code: ###  
### https://github.com/dylan-losey/me5824.git ###  
  
import matplotlib.pyplot as plt  
import torch  
import torch.nn as nn  
from torch.utils.data import Dataset, DataLoader  
import torch.optim as optim  
import numpy as np  
import random  
  
[ ]: ### Create our BC Model ###  
  
class HumanData(Dataset):  
  
    def __init__(self, data):  
        self.data = data  
  
    def __len__(self):  
        return len(self.data)  
  
    def __getitem__(self, idx):  
        return torch.FloatTensor(self.data[idx])  
  
class BC(nn.Module):  
    def __init__(self, state_dim, action_dim, hidden_dim):  
        super(BC, self).__init__()  
  
        self.state_dim = state_dim  
        self.action_dim = action_dim  
  
        self.linear1 = nn.Linear(state_dim, hidden_dim)  
        self.linear2 = nn.Linear(hidden_dim, hidden_dim)  
        self.linear3 = nn.Linear(hidden_dim, action_dim)  
  
        self.loss_func = nn.MSELoss()
```

```

def encoder(self, state):
    h1 = torch.tanh(self.linear1(state))
    h2 = torch.tanh(self.linear2(h1))
    return self.linear3(h2)

def forward(self, x):
    state = x[:, :self.state_dim]
    a_target = x[:, -self.action_dim:]
    a_predicted = self.encoder(state)
    loss = self.loss(a_predicted, a_target)
    return loss

def loss(self, a_predicted, a_target):
    return self.loss_func(a_predicted, a_target)

```

[ ]: *### Collect the human demonstrations ###*


```

N = 100                # number of demonstrations
sigma_h = 0.01         # amount of noise in the demonstration
T = 10                # each demonstration has T timesteps
D = []                # dataset of state-action pairs

for iter in range(N):
    xi = np.zeros((T, 2))
    p_robot = np.random.rand(2)
    p_goal_1 = np.random.rand(2)
    p_goal_2 = np.random.rand(2)
    choice = random.randint(0,1)
    if choice == 0:
        p_goal = p_goal_1
        target = [1, -1]
    else:
        p_goal = p_goal_2
        target = [-1, 1]

    for timestep in range(T):
        a = np.random.normal((p_goal - p_robot) / 5.0, sigma_h)
        xi[timestep, :] = np.copy(p_robot)
        D.append(p_robot.tolist() + p_goal_1.tolist() + p_goal_2.tolist() +
→target + a.tolist())
        p_robot += a
    plt.plot(p_goal[0], p_goal[1], 'ko')
    plt.plot(xi[:,0], xi[:,1], 'bo-')
plt.axis([0, 1, 0, 1])
plt.show()

```



```
[ ]: ### Train the BC Model ###

# arguments: state dimension, action dimension, hidden size
model = BC(8, 2, 32)

EPOCH = 1001
BATCH_SIZE_TRAIN = 100
LR = 0.01
LR_STEP_SIZE = 360
LR_GAMMA = 0.1

train_data = HumanData(D)
train_set = DataLoader(dataset=train_data, batch_size=BATCH_SIZE_TRAIN,
    →shuffle=True)

optimizer = optim.Adam(model.parameters(), lr=LR)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=LR_STEP_SIZE,
    →gamma=LR_GAMMA)

for epoch in range(EPOCH):
    for batch, x in enumerate(train_set):
        optimizer.zero_grad()
        loss = model(x)
        loss.backward()
        optimizer.step()
    scheduler.step()
```



```

    if epoch % 100 == 0:
        print(epoch, loss.item())
torch.save(model.state_dict(), "bc_weights")

```

```

0 0.005890688393265009
100 0.00011825011461041868
200 0.0001473638549214229
300 0.00014695516438223422
400 8.811458974378183e-05
500 0.0001227668981300667
600 9.010592475533485e-05
700 7.98568144091405e-05
800 9.984023199649528e-05
900 0.00010542211384745315
1000 9.910707012750208e-05

```

```
[ ]: ### Rollout the trained model ###
```

```

model = BC(8, 2, 32)
model.load_state_dict(torch.load("bc_weights"))

N = 10          # number of rollouts
T = 20          # each one has T timesteps
fig, ax2d = plt.subplots(5,2, figsize=(8, 20), squeeze=False)
axli = ax2d.flatten()
for iter, ax in enumerate(ax2d.flat):
    xi = np.zeros((T, 2))
    p_robot = np.random.rand(2)
    p_goal_1 = np.random.rand(2)
    p_goal_2 = np.random.rand(2)
    if iter < N/2:
        target = np.array([1, -1])
    else:
        target = np.array([-1, 1])
    for timestep in range(T):
        context = np.concatenate((p_robot, p_goal_1, p_goal_2, target))
        a = model.encoder(torch.Tensor(context)).detach().numpy()
        xi[timestep, :] = np.copy(p_robot)
        p_robot += a
    ax.plot(p_goal_1[0], p_goal_1[1], 'go', label='green fruit')
    ax.plot(p_goal_2[0], p_goal_2[1], 'bo', label='purple grapes')
    ax.plot(xi[:,0], xi[:,1], 'ko-', label='robot trajectory')
    ax.axis([0, 1, 0, 1])
    ax.legend()
    ax.set_title('Target: ' + str(target))
plt.show()

```

